

CS148_Project_1_To_Do

January 16, 2023

0.1 Introduction

Welcome to **CS148 - Data Science Fundamentals!** As we're planning to move through topics aggressively in this course, to start out, we'll look to do an end-to-end walkthrough of a datascience project, and then ask you to replicate the code yourself for a new dataset.

Please note: We don't expect you to fully grasp everything happening here in either code or theory. This content will be reviewed throughout the quarter. Rather we hope that by giving you the full perspective on a data science project it will better help to contextualize the pieces as they're covered in class

In that spirit, we will first work through an example project from end to end to give you a feel for the steps involved.

Here are the main steps:

1. Get the data
2. Visualize the data for insights
3. Preprocess the data for your machine learning algorithm
4. Select a machine learning model and train it
5. Evaluate its performance

0.2 Working with Real Data

It is best to experiment with real-data as opposed to artificial datasets.

There are many different open datasets depending on the type of problems you might be interested in!

Here are a few data repositories you could check out: - [UCI Datasets](#) - [Kaggle Datasets](#) - [AWS Datasets](#)

Below we will run through an California Housing example collected from the 1990's.

0.3 Setup

We'll start by importing a series of libraries we'll be using throughout the project.

```
[1]: import sys
      assert sys.version_info >= (3, 5) # python>=3.5
      import sklearn
      assert sklearn.__version__ >= "0.20" # sklearn >= 0.20
```

```

import numpy as np #numerical package in python
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
import plotly.io as pio
#pio.kaleido.scope.default_format = "svg"
#pio.renderers.default="notebook" DID NOT WORK

```

0.4 Intro to Data Exploration Using Pandas

In this section we will load the dataset, and visualize different features using different types of plots.

Packages we will use: - **Pandas**: is a fast, flexible and expressive data structure widely used for tabular and multidimensional datasets. - **Matplotlib**: is a 2d python plotting library which you can use to create quality figures (you can plot almost anything if you're willing to code it out!) - other plotting libraries: [seaborn](#), [ggplot2](#)

Note: If you're working in CoLab for this project, the CSV file first has to be loaded into the environment. This can be done manually using the sidebar menu option, or using the following code here.

If you're running this notebook locally on your device, simply proceed to the next step.

```

[2]: #from google.colab import files
      #files.upload()

```

We'll now begin working with Pandas. Pandas is the principle library for data management in python. It's primary mechanism of data storage is the dataframe, a two dimensional table, where each column represents a datatype, and each row a specific data element in the set.

To work with dataframes, we have to first read in the csv file and convert it to a dataframe using the code below.

```

[3]: # We'll now import the holy grail of python datascience: Pandas!
      import pandas as pd
      import os
      HOUSING_PATH = os.path.join("datasets", "housing", "housing.csv")
      housing = pd.read_csv(HOUSING_PATH)

```

```

[4]: housing.head() # show the first few elements of the dataframe
      # typically this is the first thing you do

```

```
# to see how the dataframe looks like
```

```
[4]: longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.23 37.88 41.0 880.0 129.0
1 -122.22 37.86 21.0 7099.0 1106.0
2 -122.24 37.85 52.0 1467.0 190.0
3 -122.25 37.85 52.0 1274.0 235.0
4 -122.25 37.85 52.0 1627.0 280.0

population households median_income median_house_value ocean_proximity
0 322.0 126.0 8.3252 452600.0 NEAR BAY
1 2401.0 1138.0 8.3014 358500.0 NEAR BAY
2 496.0 177.0 7.2574 352100.0 NEAR BAY
3 558.0 219.0 5.6431 341300.0 NEAR BAY
4 565.0 259.0 3.8462 342200.0 NEAR BAY
```

A dataset may have different types of features - real valued - Discrete (integers) - categorical (strings) - Boolean

The two categorical features are essentially the same as you can always map a categorical string/character to an integer.

In the dataset example, all our features are real valued floats, except ocean proximity which is categorical.

```
[5]: # to see a concise summary of data types, null values, and counts
# use the info() method on the dataframe
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[6]: # you can access individual columns similarly
# to accessing elements in a python dict
```

```
housing["ocean_proximity"].head() # added head() to avoid printing many columns.
↩.
```

```
[6]: 0    NEAR BAY
      1    NEAR BAY
      2    NEAR BAY
      3    NEAR BAY
      4    NEAR BAY
      Name: ocean_proximity, dtype: object
```

```
[7]: # to access a particular row we can use iloc
      housing.iloc[1]
```

```
[7]: longitude          -122.22
      latitude           37.86
      housing_median_age    21.0
      total_rooms          7099.0
      total_bedrooms       1106.0
      population           2401.0
      households           1138.0
      median_income         8.3014
      median_house_value   358500.0
      ocean_proximity      NEAR BAY
      Name: 1, dtype: object
```

```
[8]: # one other function that might be useful is
      # value_counts(), which counts the number of occurrences
      # for categorical features
      housing["ocean_proximity"].value_counts()
```

```
[8]: <1H OCEAN      9136
      INLAND       6551
      NEAR OCEAN   2658
      NEAR BAY     2290
      ISLAND        5
      Name: ocean_proximity, dtype: int64
```

```
[9]: # The describe function compiles your typical statistics for each
      # column
      housing.describe()
```

```
[9]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	

75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

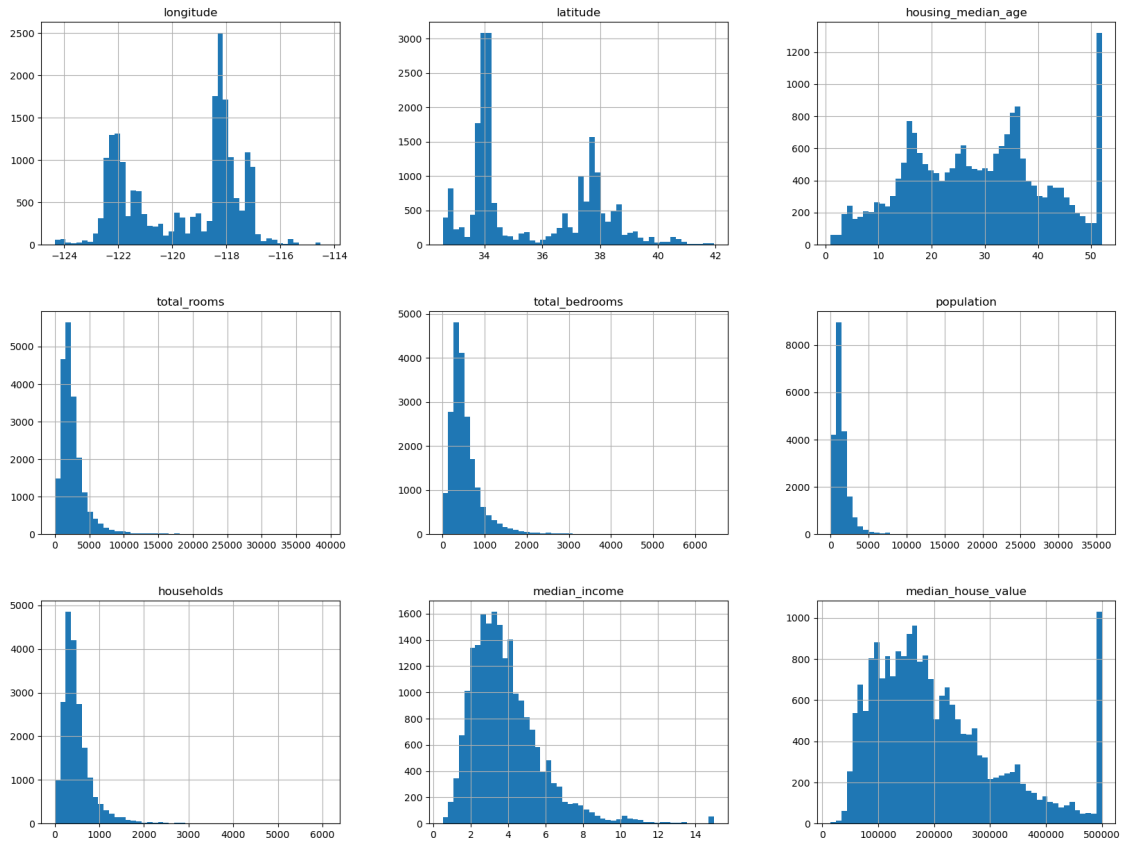
	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000


If you want to learn about different ways of accessing elements or other functions it's useful to check out the getting started section [here](#)

0.5 Let's start visualizing the dataset

```
[10]: # We can draw a histogram for each of the dataframes features
      # using the hist function
housing.hist(bins=50, figsize=(20,15))
      # save_fig("attribute_histogram_plots")
plt.show() # pandas internally uses matplotlib, and to display all the figures
           # the show() function must be called
```



```
[11]: # if you want to have a histogram on an individual feature:
housing["median_income"].hist()
plt.show()
```



We can convert a floating point feature to a categorical feature by binning or by defining a set of intervals.


For example, to bin the households based on median_income we can use the `pd.cut` function

```
[12]: # assign each bin a categorical value [1, 2, 3, 4, 5] in this case.
housing["income_cat"] = pd.cut(housing["median_income"],
                                bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                                labels=[1, 2, 3, 4, 5])

housing["income_cat"].value_counts()
```

```
[12]: 3    7236
      2    6581
      4    3639
      5    2362
      1     822
      Name: income_cat, dtype: int64
```


```
[13]: housing["income_cat"].hist()
      plt.show()
```



Next let's visualize the household incomes based on latitude & longitude coordinates


```
[14]: ## here's a not so interesting way plotting it  
housing.plot(kind="scatter", x="longitude", y="latitude")
```

```
[14]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```

```
[15]: # we can make it look a bit nicer by using the alpha parameter,  
# it simply plots less dense areas lighter.  
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
```

```
[15]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```



```
[16]: # A more interesting plot is to color code (heatmap) the dots
# based on income. The code below achieves this

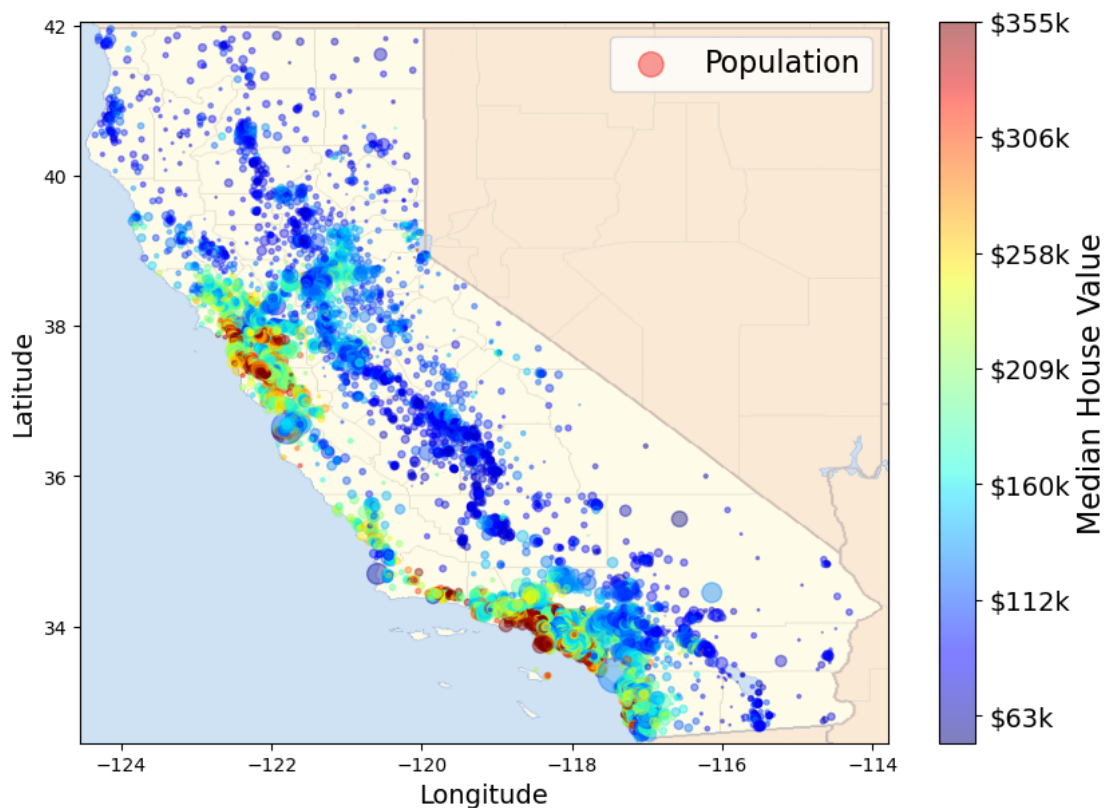
# Please note: In order for this to work, ensure that you've loaded an image
# of california (california.png) into this directory prior to running this

import matplotlib.image as mpimg
CALI_IMAGE_PATH = os.path.join("images", "california.png")
california_img = mpimg.imread(CALI_IMAGE_PATH)
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=housing['population']/100, label="Population",
                  c="median_house_value", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )
# overlay the california map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)
```

```
# setting up heatmap colors based on median_house_value feature

prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cb = plt.colorbar()
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
    ↪fontsize=14)
cb.set_label('Median House Value', fontsize=16)

plt.legend(fontsize=16)
plt.show()
```



Not surprisingly, the most expensive houses are concentrated around the San Francisco/Los Angeles areas.

Up until now we have only visualized feature histograms and basic statistics.

When developing machine learning models the predictiveness of a feature for a particular target of interest is what's important.

It may be that only a few features are useful for the target at hand, or features may need to be augmented by applying certain transformations.

None the less we can explore this using correlation matrices.

```
[17]: corr_matrix = housing.corr()
```

```
[18]: corr_matrix["median_house_value"]
```

```
[18]: longitude          -0.045967
latitude          -0.144160
housing_median_age    0.105623
total_rooms          0.134153
total_bedrooms        0.049686
population          -0.024650
households           0.065843
median_income         0.688075
median_house_value    1.000000
Name: median_house_value, dtype: float64
```

```
[19]: # for example if the target is "median_house_value", most correlated features
      ↳ can be sorted
      # which happens to be "median_income". This also intuitively makes sense.
      corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[19]: median_house_value    1.000000
      median_income        0.688075
      total_rooms          0.134153
      housing_median_age    0.105623
      households           0.065843
      total_bedrooms        0.049686
      population          -0.024650
      longitude          -0.045967
      latitude          -0.144160
      Name: median_house_value, dtype: float64
```

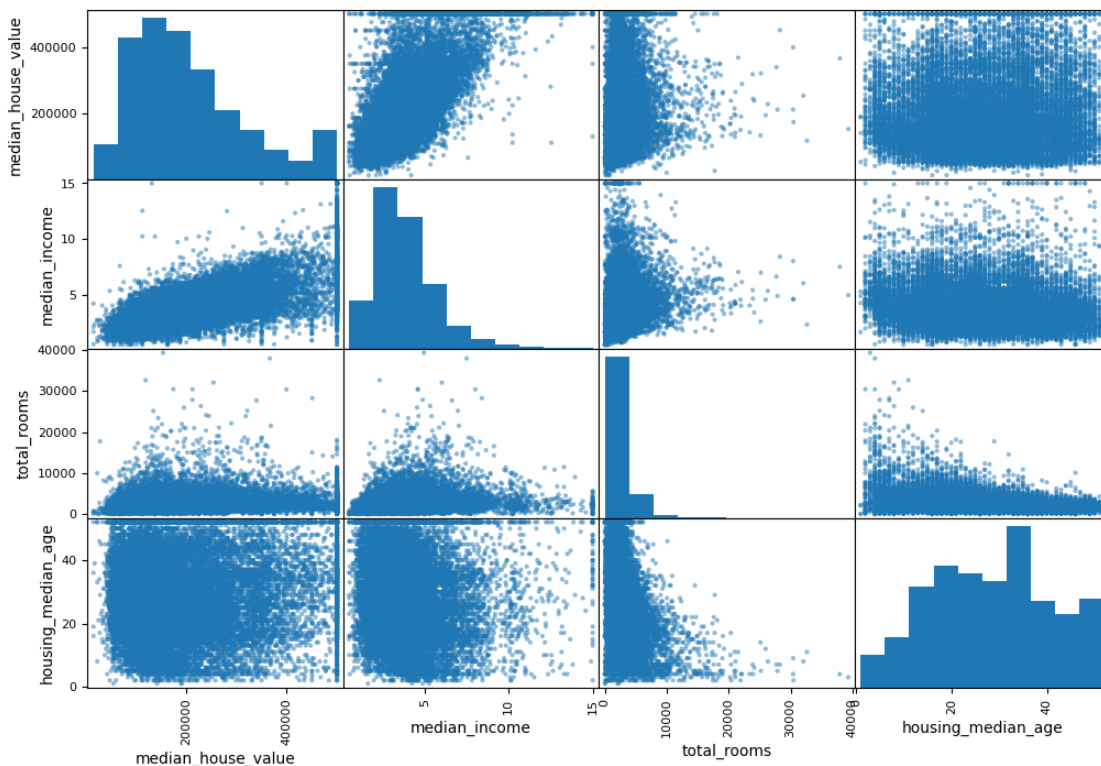
```
[20]: # the correlation matrix for different attributes/features can also be plotted
      # some features may show a positive correlation/negative correlation or
      # it may turn out to be completely random!
      from pandas.plotting import scatter_matrix
      attributes = ["median_house_value", "median_income", "total_rooms",
                    "housing_median_age"]
      scatter_matrix(housing[attributes], figsize=(12, 8))
```

```
[20]: array([[<AxesSubplot:xlabel='median_house_value', ylabel='median_house_value'>,
      <AxesSubplot:xlabel='median_income', ylabel='median_house_value'>,
      <AxesSubplot:xlabel='total_rooms', ylabel='median_house_value'>,
      <AxesSubplot:xlabel='housing_median_age', ylabel='median_house_value'>],
      [<AxesSubplot:xlabel='median_house_value', ylabel='median_income'>,
      <AxesSubplot:xlabel='median_income', ylabel='median_income'>,
      <AxesSubplot:xlabel='total_rooms', ylabel='median_income'>],
```

```

    <AxesSubplot:xlabel='housing_median_age', ylabel='median_income'>],
    [<AxesSubplot:xlabel='median_house_value', ylabel='total_rooms'>,
    <AxesSubplot:xlabel='median_income', ylabel='total_rooms'>,
    <AxesSubplot:xlabel='total_rooms', ylabel='total_rooms'>,
    <AxesSubplot:xlabel='housing_median_age', ylabel='total_rooms'>],
    [<AxesSubplot:xlabel='median_house_value', ylabel='housing_median_age'>,
    <AxesSubplot:xlabel='median_income', ylabel='housing_median_age'>,
    <AxesSubplot:xlabel='total_rooms', ylabel='housing_median_age'>,
    <AxesSubplot:xlabel='housing_median_age',
    ylabel='housing_median_age'>]],
    dtype=object)

```



```


[21]: # median income vs median house value plot plot 2 in the first row of top figure
housing.plot(kind="scatter", x="median_income", y="median_house_value",
            alpha=0.1)
plt.axis([0, 16, 0, 550000])

```

```

[21]: (0.0, 16.0, 0.0, 550000.0)


```



```
[22]: # obtain new correlations
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
[22]: median_house_value    1.000000
      median_income       0.688075
      total_rooms        0.134153
      housing_median_age  0.105623
      households         0.065843
      total_bedrooms     0.049686
      population        -0.024650
      longitude         -0.045967
      latitude          -0.144160
      Name: median_house_value, dtype: float64
```

```
[23]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```



0.6 Preparing Dataset for ML

0.6.1 Augmenting Features

New features can be created by combining different columns from our data set.

- $\text{rooms_per_household} = \text{total_rooms} / \text{households}$
- $\text{bedrooms_per_room} = \text{total_bedrooms} / \text{total_rooms}$
- etc.

```
[24]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"]=housing["population"]/housing["households"]
```

0.6.2 Dealing With Incomplete Data

```
[25]: # have you noticed when looking at the dataframe summary certain rows
# contained null values? we can't just leave them as nulls and expect our
# model to handle them for us...
sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
sample_incomplete_rows
```

```
[25]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
290	-122.16	37.77	47.0	1256.0	NaN	
341	-122.17	37.75	38.0	992.0	NaN	
538	-122.28	37.78	29.0	5154.0	NaN	
563	-122.24	37.75	45.0	891.0	NaN	
696	-122.10	37.69	41.0	746.0	NaN	

	population	households	median_income	median_house_value	\
290	570.0	218.0	4.3750	161900.0	
341	732.0	259.0	1.6196	85100.0	
538	3741.0	1273.0	2.5762	173400.0	
563	384.0	146.0	4.9489	247100.0	
696	387.0	161.0	3.9063	178400.0	

	ocean_proximity	income_cat	rooms_per_household	bedrooms_per_room	\
290	NEAR BAY	3	5.761468	NaN	
341	NEAR BAY	2	3.830116	NaN	
538	NEAR BAY	2	4.048704	NaN	
563	NEAR BAY	4	6.102740	NaN	
696	NEAR BAY	3	4.633540	NaN	

	population_per_household
290	2.614679
341	2.826255
538	2.938727
563	2.630137
696	2.403727

```
[26]: sample_incomplete_rows.dropna(subset=["total_bedrooms"]) # option 1: simply
      ↪ drop rows that have null values
```

```
[26]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms,
population, households, median_income, median_house_value, ocean_proximity,
income_cat, rooms_per_household, bedrooms_per_room, population_per_household]
Index: []
```

```
[27]: sample_incomplete_rows.drop("total_bedrooms", axis=1) # option 2: drop
      ↪ the complete feature
```

```
[27]:
```

	longitude	latitude	housing_median_age	total_rooms	population	\
290	-122.16	37.77	47.0	1256.0	570.0	
341	-122.17	37.75	38.0	992.0	732.0	
538	-122.28	37.78	29.0	5154.0	3741.0	
563	-122.24	37.75	45.0	891.0	384.0	
696	-122.10	37.69	41.0	746.0	387.0	

	households	median_income	median_house_value	ocean_proximity	income_cat	\
290	218.0	4.3750	161900.0	NEAR BAY	3	
341	259.0	1.6196	85100.0	NEAR BAY	2	
538	1273.0	2.5762	173400.0	NEAR BAY	2	
563	146.0	4.9489	247100.0	NEAR BAY	4	
696	161.0	3.9063	178400.0	NEAR BAY	3	

	rooms_per_household	bedrooms_per_room	population_per_household
290	5.761468	NaN	2.614679
341	3.830116	NaN	2.826255
538	4.048704	NaN	2.938727
563	6.102740	NaN	2.630137
696	4.633540	NaN	2.403727

```
[28]: median = housing["total_bedrooms"].median()
sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option_
↳3: replace na values with median values
sample_incomplete_rows
```

```
[28]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
290	-122.16	37.77	47.0	1256.0	435.0	
341	-122.17	37.75	38.0	992.0	435.0	
538	-122.28	37.78	29.0	5154.0	435.0	
563	-122.24	37.75	45.0	891.0	435.0	
696	-122.10	37.69	41.0	746.0	435.0	

	population	households	median_income	median_house_value	\
290	570.0	218.0	4.3750	161900.0	
341	732.0	259.0	1.6196	85100.0	
538	3741.0	1273.0	2.5762	173400.0	
563	384.0	146.0	4.9489	247100.0	
696	387.0	161.0	3.9063	178400.0	

	ocean_proximity	income_cat	rooms_per_household	bedrooms_per_room	\
290	NEAR BAY	3	5.761468	NaN	
341	NEAR BAY	2	3.830116	NaN	
538	NEAR BAY	2	4.048704	NaN	
563	NEAR BAY	4	6.102740	NaN	
696	NEAR BAY	3	4.633540	NaN	

	population_per_household
290	2.614679
341	2.826255
538	2.938727
563	2.630137
696	2.403727

Now that we've played around with this, lets finalize this approach by replacing the nulls in our

final dataset

```
[29]: housing["total_bedrooms"].fillna(median, inplace=True)
```

Could you think of another plausible imputation for this dataset?

```
[30]: median = housing["bedrooms_per_room"].median()
housing["bedrooms_per_room"].fillna(median, inplace=True)
```

0.6.3 Dealing with Non-Numeric Data

So we're almost ready to feed our dataset into a machine learning model, but we're not quite there yet!

Generally speaking all models can only work with numeric data, which means that if you have Categorical data you want included in your model, you'll need to do a numeric conversion. We'll explore this more later, but for now we'll take one approach to converting our `ocean_proximity` field into a numeric one.

```
[31]: from sklearn.preprocessing import LabelEncoder

# creating instance of labelencoder
labelencoder = LabelEncoder()
# Assigning numerical values and storing in another column
housing['ocean_proximity'] = labelencoder.
    ↪fit_transform(housing['ocean_proximity'])
housing.head()
```

```
[31]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity	\
0	322.0	126.0	8.3252	452600.0	3	
1	2401.0	1138.0	8.3014	358500.0	3	
2	496.0	177.0	7.2574	352100.0	3	
3	558.0	219.0	5.6431	341300.0	3	
4	565.0	259.0	3.8462	342200.0	3	

	income_cat	rooms_per_household	bedrooms_per_room	population_per_household
0	5	6.984127	0.146591	2.555556
1	5	6.238137	0.155797	2.109842
2	5	8.288136	0.129516	2.802260
3	4	5.817352	0.184458	2.547945
4	3	6.281853	0.172096	2.181467

0.6.4 Divide up the Dataset for Machine Learning

After having cleaned your dataset you're ready to train your machine learning model.

To do so you'll aim to divide your data into: - train set - test set

In some cases you might also have a validation set as well for tuning hyperparameters (don't worry if you're not familiar with this term yet..)

In supervised learning setting your train set and test set should contain (**feature**, **target**) tuples. - **feature**: is the input to your model - **target**: is the ground truth label - when target is categorical the task is a classification task - when target is floating point the task is a regression task

We will make use of [scikit-learn](#) python package for preprocessing.

Scikit learn is pretty well documented and if you get confused at any point simply look up the function/object!

```
[32]: from sklearn.model_selection import StratifiedShuffleSplit
      # let's first start by creating our train and test sets
      split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
      for train_index, test_index in split.split(housing, housing["income_cat"]):
          train_set = housing.loc[train_index]
          test_set = housing.loc[test_index]

[33]: housing_training = train_set.drop("median_house_value", axis=1) # drop labels
      # for training set features
      # the input to the model
      # should not contain the true label
      housing_labels = train_set["median_house_value"].copy()

[34]: housing_testing = test_set.drop("median_house_value", axis=1) # drop labels for
      # training set features
      # the input to the model
      # should not contain the true label
      housing__test_labels = test_set["median_house_value"].copy()
```

0.6.5 Select a model and train

Once we have prepared the dataset it's time to choose a model.

As our task is to predict the median_house_value (a floating value), regression is well suited for this.

```
[35]: from sklearn.linear_model import LinearRegression

      lin_reg = LinearRegression()
      lin_reg.fit(housing_training, housing_labels)

[35]: LinearRegression()
```

```
[36]: # let's try our model on a few testing instances
data = housing_testing.iloc[:5]
labels = housing__test_labels.iloc[:5]

print("Predictions:", lin_reg.predict(data))
print("Actual labels:", list(labels))
```

```
Predictions: [423969.88487578 298939.35075414 227783.64452614 185030.00931653
244874.28030822]
```

```
Actual labels: [500001.0, 162500.0, 204600.0, 159700.0, 184000.0]
```

We can evaluate our model using certain metrics, a fitting metric for regresison is the mean-squared-loss

$$L(\hat{Y}, Y) = \sum_i^N (\hat{y}_i - y_i)^2$$

where \hat{y} is the predicted value, and y is the ground truth label.

```
[37]: from sklearn.metrics import mean_squared_error

preds = lin_reg.predict(housing_testing)
mse = mean_squared_error(housing__test_labels, preds)
rmse = np.sqrt(mse)
rmse
```

```
[37]: 67538.56731651393
```

Is this a good result? What do you think an acceptable error rate is for this sort of problem?

1 TODO: Applying the end-end ML steps to a different dataset.

Ok now it's time to get to work! We will apply what we've learnt to another dataset (airbnb dataset). For this project we will attempt to **predict the airbnb rental price based on other features in our given dataset.**

2 Visualizing Data

2.0.1 Load the data + statistics

Let's do the following set of tasks to get us warmed up: - load the dataset - display the first few rows of the data - drop the following columns: name, host_id, host_name, last_review, neighbourhood - display a summary of the statistics of the loaded data

```
[38]: import pandas as pd
import os
import sys
assert sys.version_info >= (3, 5) # python>=3.5
```

```

import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt
import plotly.express as px

```

```

[39]: # NOTE TO GRADER, I WAS NOT ABLE TO PROPERLY EXPORT PLOTLY GRAPHS TO PDF, SO I
      ↪HAVE ADDED THEIR PNGS AT THE END OF THE DOCUMENT.
# NOTE TO GRADER, I WAS NOT ABLE TO PROPERLY EXPORT PLOTLY GRAPHS TO PDF, SO I
      ↪HAVE ADDED THEIR PNGS AT THE END OF THE DOCUMENT.
# NOTE TO GRADER, I WAS NOT ABLE TO PROPERLY EXPORT PLOTLY GRAPHS TO PDF, SO I
      ↪HAVE ADDED THEIR PNGS AT THE END OF THE DOCUMENT.
# NOTE TO GRADER, I WAS NOT ABLE TO PROPERLY EXPORT PLOTLY GRAPHS TO PDF, SO I
      ↪HAVE ADDED THEIR PNGS AT THE END OF THE DOCUMENT.
# NOTE TO GRADER, I WAS NOT ABLE TO PROPERLY EXPORT PLOTLY GRAPHS TO PDF, SO I
      ↪HAVE ADDED THEIR PNGS AT THE END OF THE DOCUMENT.
# NOTE TO GRADER, I WAS NOT ABLE TO PROPERLY EXPORT PLOTLY GRAPHS TO PDF, SO I
      ↪HAVE ADDED THEIR PNGS AT THE END OF THE DOCUMENT.

```

```

[40]: #1 Load the dataset
AIRBNB_PATH = os.path.join("datasets","airbnb","AB_NYC_2019.csv")
airbnb = pd.read_csv(AIRBNB_PATH)

```

```

[41]: #2 Display the first few rows of the data
airbnb.head()

```

```

[41]:      id      name  host_id \
0  2539  Clean & quiet apt home by the park    2787
1  2595                Skylit Midtown Castle    2845
2  3647  THE VILLAGE OF HARLEM...NEW YORK !    4632
3  3831      Cozy Entire Floor of Brownstone    4869
4  5022  Entire Apt: Spacious Studio/Loft by central park    7192

      host_name  neighbourhood_group  neighbourhood  latitude  longitude \
0         John          Brooklyn    Kensington  40.64749  -73.97237
1      Jennifer          Manhattan      Midtown  40.75362  -73.98377
2     Elisabeth          Manhattan      Harlem  40.80902  -73.94190
3  LisaRoxanne          Brooklyn  Clinton Hill  40.68514  -73.95976
4         Laura          Manhattan    East Harlem  40.79851  -73.94399

```

	room_type	price	minimum_nights	number_of_reviews	last_review	\
0	Private room	149	1	9	2018-10-19	
1	Entire home/apt	225	1	45	2019-05-21	
2	Private room	150	3	0	NaN	
3	Entire home/apt	89	1	270	2019-07-05	
4	Entire home/apt	80	10	9	2018-11-19	

	reviews_per_month	calculated_host_listings_count	availability_365
0	0.21	6	365
1	0.38	2	355
2	NaN	1	365
3	4.64	1	194
4	0.10	1	0

```
[42]: #3 Drop the following columns:name, host_id, host_name, last_review,
      ↪neighbourhood
airbnb_modified = airbnb.drop(["name", "host_id", "host_name", "last_review"],
      ↪axis=1)
```

```
[43]: #4 Display a summary of the statistics of the loaded data
airbnb_modified.describe()
```

```
[43]:
```

	id	latitude	longitude	price	minimum_nights	\
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000	
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962	
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550	
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000	
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000	
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000	
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000	
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000	

	number_of_reviews	reviews_per_month	calculated_host_listings_count	\
count	48895.000000	38843.000000	48895.000000	
mean	23.274466	1.373221	7.143982	
std	44.550582	1.680442	32.952519	
min	0.000000	0.010000	1.000000	
25%	1.000000	0.190000	1.000000	
50%	5.000000	0.720000	1.000000	
75%	24.000000	2.020000	2.000000	
max	629.000000	58.500000	327.000000	

	availability_365
count	48895.000000
mean	112.781327
std	131.622289

```

min            0.000000
25%            0.000000
50%            45.000000
75%            227.000000
max            365.000000

```

2.0.2 Some Basic Visualizations

Let's try another popular python graphics library: Plotly.

You can find documentation and all the examples you'll need here: [Plotly Documentation](#)

Let's start out by getting a better feel for the distribution of rentals in the market.

####Generate a pie chart showing the distribution of rental units across NYC's 5 Boroughs (neighbourhood_groups in the dataset)####

```
[44]: airbnb_modified["neighbourhood_group"].value_counts()
```

```

[44]: Manhattan      21661
      Brooklyn      20104
      Queens         5666
      Bronx          1091
      Staten Island   373
      Name: neighbourhood_group, dtype: int64

```

```

[45]: #Pie chart showing the distribution of rental units across NYC's 5 Boroughs
      fig = px.pie(airbnb_modified, names='neighbourhood_group', title='The
      ↪distribution of rental units across NYC\'s 5 Boroughs')
      fig.show()

```

Plot the total number_of_reviews per neighbourhood_group We now want to see the total number of reviews left for each neighborhood group in the form of a Bar Chart (where the X-axis is the neighbourhood group and the Y-axis is a count of review.

This is a two step process: 1. You'll have to sum up the reviews per neighbourhood group (**hint! try using the groupby function**) 2. Then use Plotly to generate the graph

```

[46]: #1 Sum up reviews
      bar_data = airbnb_modified.groupby("neighbourhood_group")["number_of_reviews"].
      ↪count()

```

```

[47]: #2 Plot the bar graph
      fig = px.bar(bar_data, y = "number_of_reviews", title = " Number of reviews per
      ↪neighbourhood_group")
      fig.show()

```

2.0.3 Plot a map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

For reference you can use the Matplotlib code above to replicate this graph here.

```
[48]: sampled_airbnb = airbnb_modified.sample(frac = 0.1) # shuffling data and taking a
      ↪ subset of the data

[49]: # Cutting of the price at value 500 for better map plot.
      sampled_airbnb["price"][sampled_airbnb["price"] >= 500] = 500


[50]: import matplotlib.image as mpimg
      NYC_IMAGE_PATH = os.path.join("images", "nyc.png")
      nyc_img = mpimg.imread(NYC_IMAGE_PATH)

      ax = sampled_airbnb.plot(kind="scatter", x="longitude", y="latitude",
                               figsize=(10,7), label="Price", s = 5,
                               cmap=plt.get_cmap("jet"), colorbar=False, c = "price",
                               alpha=0.3)

      # overlay the nyc map on the plotted scatter plot
      # note: plt.imshow still refers to the most recent figure
      # that hasn't been plotted yet.
      summary_airbnb = airbnb_modified.describe()
      lon_min = summary_airbnb["longitude"]["min"]
      lon_max = summary_airbnb["longitude"]["max"]
      lat_min = summary_airbnb["latitude"]["min"]
      lat_max = summary_airbnb["latitude"]["max"]
      plt.imshow(nyc_img, extent=[lon_min, lon_max, lat_min, lat_max], alpha = 0.5,
                 cmap=plt.get_cmap("jet"))
      plt.ylabel("Latitude", fontsize=14)
      plt.xlabel("Longitude", fontsize=14)
      plt.title("Map of airbnbs throughout New York")

      prices = sampled_airbnb["price"]
      tick_values = np.linspace(prices.min(), prices.max(), 11)
      cb = plt.colorbar()
      cb.ax.set_yticklabels(["$%d"%v for v in tick_values], fontsize=14)
      cb.set_label("Price per Night", fontsize=16)

      plt.show()
```

Now try to recreate this plot using Plotly's Scatterplot functionality. Note that the increased interactivity of the plot allows for some very cool functionality

```
[51]: from PIL import Image
fig = px.scatter(sampled_airbnb, x="longitude", y="latitude", color = "price",
               ↪title = "Map of Airbnbs throughout New York")
import base64
#set a local image as a background
plotly_logo = base64.b64encode(open(NYC_IMAGE_PATH, 'rb').read())

fig.update_layout(
    images= [dict(
        source='data:image/png;base64,{}'.format(plotly_logo.
        ↪decode()),
        xref="paper", yref="paper",
        x=0, y=1,
        sizex=1, sizey=1,
        xanchor="left",
        yanchor="top",
        sizing="stretch",
        layer="below")])

summary_airbnb = airbnb_modified.describe()
```

```

lon_min = summary_airbnb["longitude"]["min"]
lon_max = summary_airbnb["longitude"]["max"]
lat_min = summary_airbnb["latitude"]["min"]
lat_max = summary_airbnb["latitude"]["max"]
fig.update_xaxes(range=[lon_min, lon_max])
fig.update_yaxes(range=[lat_min, lat_max])

fig.show()

```

2.0.4 Use Plotly to plot the average price of room types in Brooklyn who have at least 10 Reviews.

Like with the previous example you'll have to do a little bit of data engineering before you actually generate the plot.

Generally I'd recommend the following series of steps: 1. Filter the data by neighborhood group and number of reviews to arrive at the subset of data relevant to this graph. 2. Groupby the room type 3. Take the mean of the price for each roomtype group 4. FINALLY (seriously!?!?) plot the result

```

[52]: # Filter the data by neighborhood group and number of reviews to arrive at the
      ↪ subset of data relevant to this graph.
brooklyn_airbnb = airbnb_modified.loc[airbnb_modified["neighbourhood_group"] ==
      ↪ "Brooklyn"]
at_least_10_brooklyn_airbnb = brooklyn_airbnb.
      ↪ loc[brooklyn_airbnb["number_of_reviews"] >= 10]

```

```

[53]: #Groupby the room type
grouped_data = at_least_10_brooklyn_airbnb[["room_type", "price"]].
      ↪ groupby("room_type")

```

```

[54]: #take the mean
averaged_grouped_data = grouped_data.mean()

```

```

[55]: #plot the result
fig = px.bar(averaged_grouped_data, title = "Average price of rooms in Brooklyn
      ↪ who have at least 10 Reviews")
fig.show()

```

3 Prepare the Data

3.0.1 Feature Engineering

Let's create a new binned feature, `price_cat` that will divide our dataset into quintiles (1-5) in terms of price level (you can choose the levels to assign)

Do a value count to check the distribution of values

```
[56]: # assign each bin a categorical value [1, 2, 3, 4, 5] in this case.
airbnb_modified["price_cat"] = pd.cut(airbnb_modified["price"],
                                     bins=[-1, 50, 100, 250, 1000, 20000],
                                     labels=[1, 2, 3, 4, 5])

airbnb_modified["price_cat"].value_counts()
```

```
[56]: 3    19759
      2    17367
      1     6561
      4     4969
      5      239
      Name: price_cat, dtype: int64
```

Now engineer at least one new feature.

```
[57]: # Computing maximum number of stays from availability and minimum nights can be
      ↪useful.
airbnb_modified["maximum_stays"] = airbnb_modified["availability_365"] /
      ↪airbnb_modified["minimum_nights"]
```

```
[58]: airbnb_modified.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     48895 non-null  int64
1   neighbourhood_group                   48895 non-null  object
2   neighbourhood                         48895 non-null  object
3   latitude                             48895 non-null  float64
4   longitude                             48895 non-null  float64
5   room_type                             48895 non-null  object
6   price                                 48895 non-null  int64
7   minimum_nights                       48895 non-null  int64
8   number_of_reviews                    48895 non-null  int64
9   reviews_per_month                    38843 non-null  float64
10  calculated_host_listings_count        48895 non-null  int64
11  availability_365                      48895 non-null  int64
12  price_cat                             48895 non-null  category
13  maximum_stays                        48895 non-null  float64
dtypes: category(1), float64(4), int64(6), object(3)
memory usage: 4.9+ MB
```

3.0.2 Data Imputation

Determine if there are any null-values and if there are impute them.

```
[59]: #reviews_per_month feature has null values only when number_of_reviews are 0,
      ↪so it is wise to fill reviews_per_month with 0.
airbnb_modified[airbnb_modified.isnull().any(axis=1)].head()
```

```
[59]:
```

	id	neighbourhood_group	neighbourhood	latitude	longitude	\
2	3647	Manhattan	Harlem	40.80902	-73.94190	
19	7750	Manhattan	East Harlem	40.79685	-73.94872	
26	8700	Manhattan	Inwood	40.86754	-73.92639	
36	11452	Brooklyn	Bedford-Stuyvesant	40.68876	-73.94312	
38	11943	Brooklyn	Flatbush	40.63702	-73.96327	

	room_type	price	minimum_nights	number_of_reviews	\
2	Private room	150	3	0	
19	Entire home/apt	190	7	0	
26	Private room	80	4	0	
36	Private room	35	60	0	
38	Private room	150	1	0	

	reviews_per_month	calculated_host_listings_count	availability_365	\
2	NaN	1	365	
19	NaN	2	249	
26	NaN	1	0	
36	NaN	1	365	
38	NaN	1	365	

	price_cat	maximum_stays
2	3	121.666667
19	3	35.571429
26	2	0.000000
36	1	6.083333
38	3	365.000000

```
[60]: airbnb_modified["reviews_per_month"].fillna(0,inplace = True)
```

```
[61]: airbnb_modified[airbnb_modified.isnull().any(axis=1)].head() #there is no null
      ↪value left.
```

```
[61]: Empty DataFrame
Columns: [id, neighbourhood_group, neighbourhood, latitude, longitude,
room_type, price, minimum_nights, number_of_reviews, reviews_per_month,
calculated_host_listings_count, availability_365, price_cat, maximum_stays]
Index: []
```

3.0.3 Numeric Conversions

Finally, review what features in your dataset are non-numeric and convert them.

```
[62]: airbnb_modified.info()
# Non-numeric features
# 1) neighbourhood_group
# 2) neighbourhood
# 3) room_type
# 4) price_cat # new feature created by binning prices

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     48895 non-null  int64
1   neighbourhood_group                   48895 non-null  object
2   neighbourhood                         48895 non-null  object
3   latitude                             48895 non-null  float64
4   longitude                             48895 non-null  float64
5   room_type                             48895 non-null  object
6   price                                 48895 non-null  int64
7   minimum_nights                       48895 non-null  int64
8   number_of_reviews                    48895 non-null  int64
9   reviews_per_month                    48895 non-null  float64
10  calculated_host_listings_count        48895 non-null  int64
11  availability_365                      48895 non-null  int64
12  price_cat                             48895 non-null  category
13  maximum_stays                        48895 non-null  float64
dtypes: category(1), float64(4), int64(6), object(3)
memory usage: 4.9+ MB
```

```
[63]: from sklearn.preprocessing import LabelEncoder

# creating instance of labelencoder
labelencoder = LabelEncoder()
# Assigning numerical values and storing in another column
airbnb_modified['neighbourhood_group'] = labelencoder.
    ↳fit_transform(airbnb_modified['neighbourhood_group'])
airbnb_modified['neighbourhood'] = labelencoder.
    ↳fit_transform(airbnb_modified['neighbourhood'])
airbnb_modified['room_type'] = labelencoder.
    ↳fit_transform(airbnb_modified['room_type'])
airbnb_modified['price_cat'] = labelencoder.
    ↳fit_transform(airbnb_modified['price_cat'])
airbnb_modified.head()
```

```
[63]:      id  neighbourhood_group  neighbourhood  latitude  longitude  room_type  \
0   2539                    1           108   40.64749   -73.97237         1
1   2595                    2           127   40.75362   -73.98377         0
```

2	3647	2	94	40.80902	-73.94190	1
3	3831	1	41	40.68514	-73.95976	0
4	5022	2	61	40.79851	-73.94399	0

	price	minimum_nights	number_of_reviews	reviews_per_month	\
0	149	1	9	0.21	
1	225	1	45	0.38	
2	150	3	0	0.00	
3	89	1	270	4.64	
4	80	10	9	0.10	

	calculated_host_listings_count	availability_365	price_cat	maximum_stays
0	6	365	2	365.000000
1	2	355	2	355.000000
2	1	365	2	121.666667
3	1	194	1	194.000000
4	1	0	1	0.000000

4 Prepare data for Machine Learning

4.0.1 Set aside 20% of the data as test test (80% train, 20% test).

Using our StratifiedShuffleSplit function example from above, let's split our data into a 80/20 Training/Testing split using neighbourhood_group to partition the dataset

```
[64]: from sklearn.model_selection import StratifiedShuffleSplit

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(airbnb_modified,
    ↪airbnb_modified["neighbourhood_group"]):
    train_set = airbnb_modified.loc[train_index]
    test_set = airbnb_modified.loc[test_index]
```

Finally, remove your labels price from your testing and training cohorts, and create separate label features.

```
[65]: airbnb_training_data = train_set.drop("price", axis=1) # drop labels for
    ↪training set features
airbnb_training_label = train_set["price"].copy()

[66]: airbnb_test_data = test_set.drop("price", axis=1) # drop labels for test set
    ↪features
airbnb_test_label = test_set["price"].copy()
```

5 Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using **MSE**. Provide both **test** and **train set MSE values**.

```
[67]: from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_squared_error

      lin_reg = LinearRegression()
      lin_reg.fit(airbnb_training_data, airbnb_training_label)

      train_prediction = lin_reg.predict(airbnb_training_data)
      train_mse = mean_squared_error(airbnb_training_label, train_prediction)
      print("Train MSE:", train_mse)

      test_prediction = lin_reg.predict(airbnb_test_data)
      test_mse = mean_squared_error(airbnb_test_label, test_prediction)
      print("Test MSE:", test_mse)
```

Train MSE: 42298.50070176791

Test MSE: 43045.73995015371